



# Flink Checkpoint-轻量级分布式快照

唐云 · 阿里巴巴 / 高级开发工程师

Apache Flink Community China

# CONTENT

## 目录 >>

**01 /**

总览: Checkpoint与state

**02 /**

什么是state

**03 /**

如何在Flink中使用state

**04 /**

Checkpoint的执行机制

# 01

---

## Checkpoint与State

二者的关系

---

# Checkpoint与State

Checkpoint，在Flink中是一个执行操作，最终产生的结果作为分布式快照提供容错机制。

Subtasks					
Task Metrics					
Watermarks					
Accumulators					
Checkpoints					
Back Pressure					
Overview					
History					
Summary					
Configuration					
Checkpoint Counts		Triggered: 569027	In Progress: 0	Completed: 569027	Failed: 0 Restored: 0
Latest Completed Checkpoint	ID: 569027	Completion Time: 18:19:02	End to End Duration: 3ms	State Size: 9.32 KB	<a href="#">More details</a>
Latest Failed Checkpoint	None				
Latest Savepoint	None				
Latest Restore	None				



# Checkpoint与State

State是构成checkpoint的数据构成

Details for Checkpoint 569116									
ID	Status	Acknowledged	Trigger Time	Latest Acknowledgement	End to End Duration	State Size	Buffered During Alignment	Discarded	Path
569116	✓ Completed	2/2 (100%)	18:22:00	18:22:00	2ms	9.17 KB	0 B	Yes	<checkpoint-not-externally-addressable>
Operators									
Name		Acknowledged	Latest Acknowledgment		End to End Duration	State Size	Buffered During Alignment		
Source: Custom Source		1/1 (100%)	18:22:00		2ms	0 B	0 B	Show Subtasks ▼	
Flat Map -> Sink: Print to Std. Out		1/1 (100%)	18:22:00		2ms	9.17 KB	0 B	Show Subtasks ▼	



# 02

---

什么是state

---

# 什么是state

```
env.socketTextStream("localhost",9000)
// split up the lines in pairs (2-tuples) containing: (word,1)
.flatMap(new Tokenizer())
// group by the tuple field "0" and sum up tuple field "1"
.keyBy(0).sum(1)
.print();
```

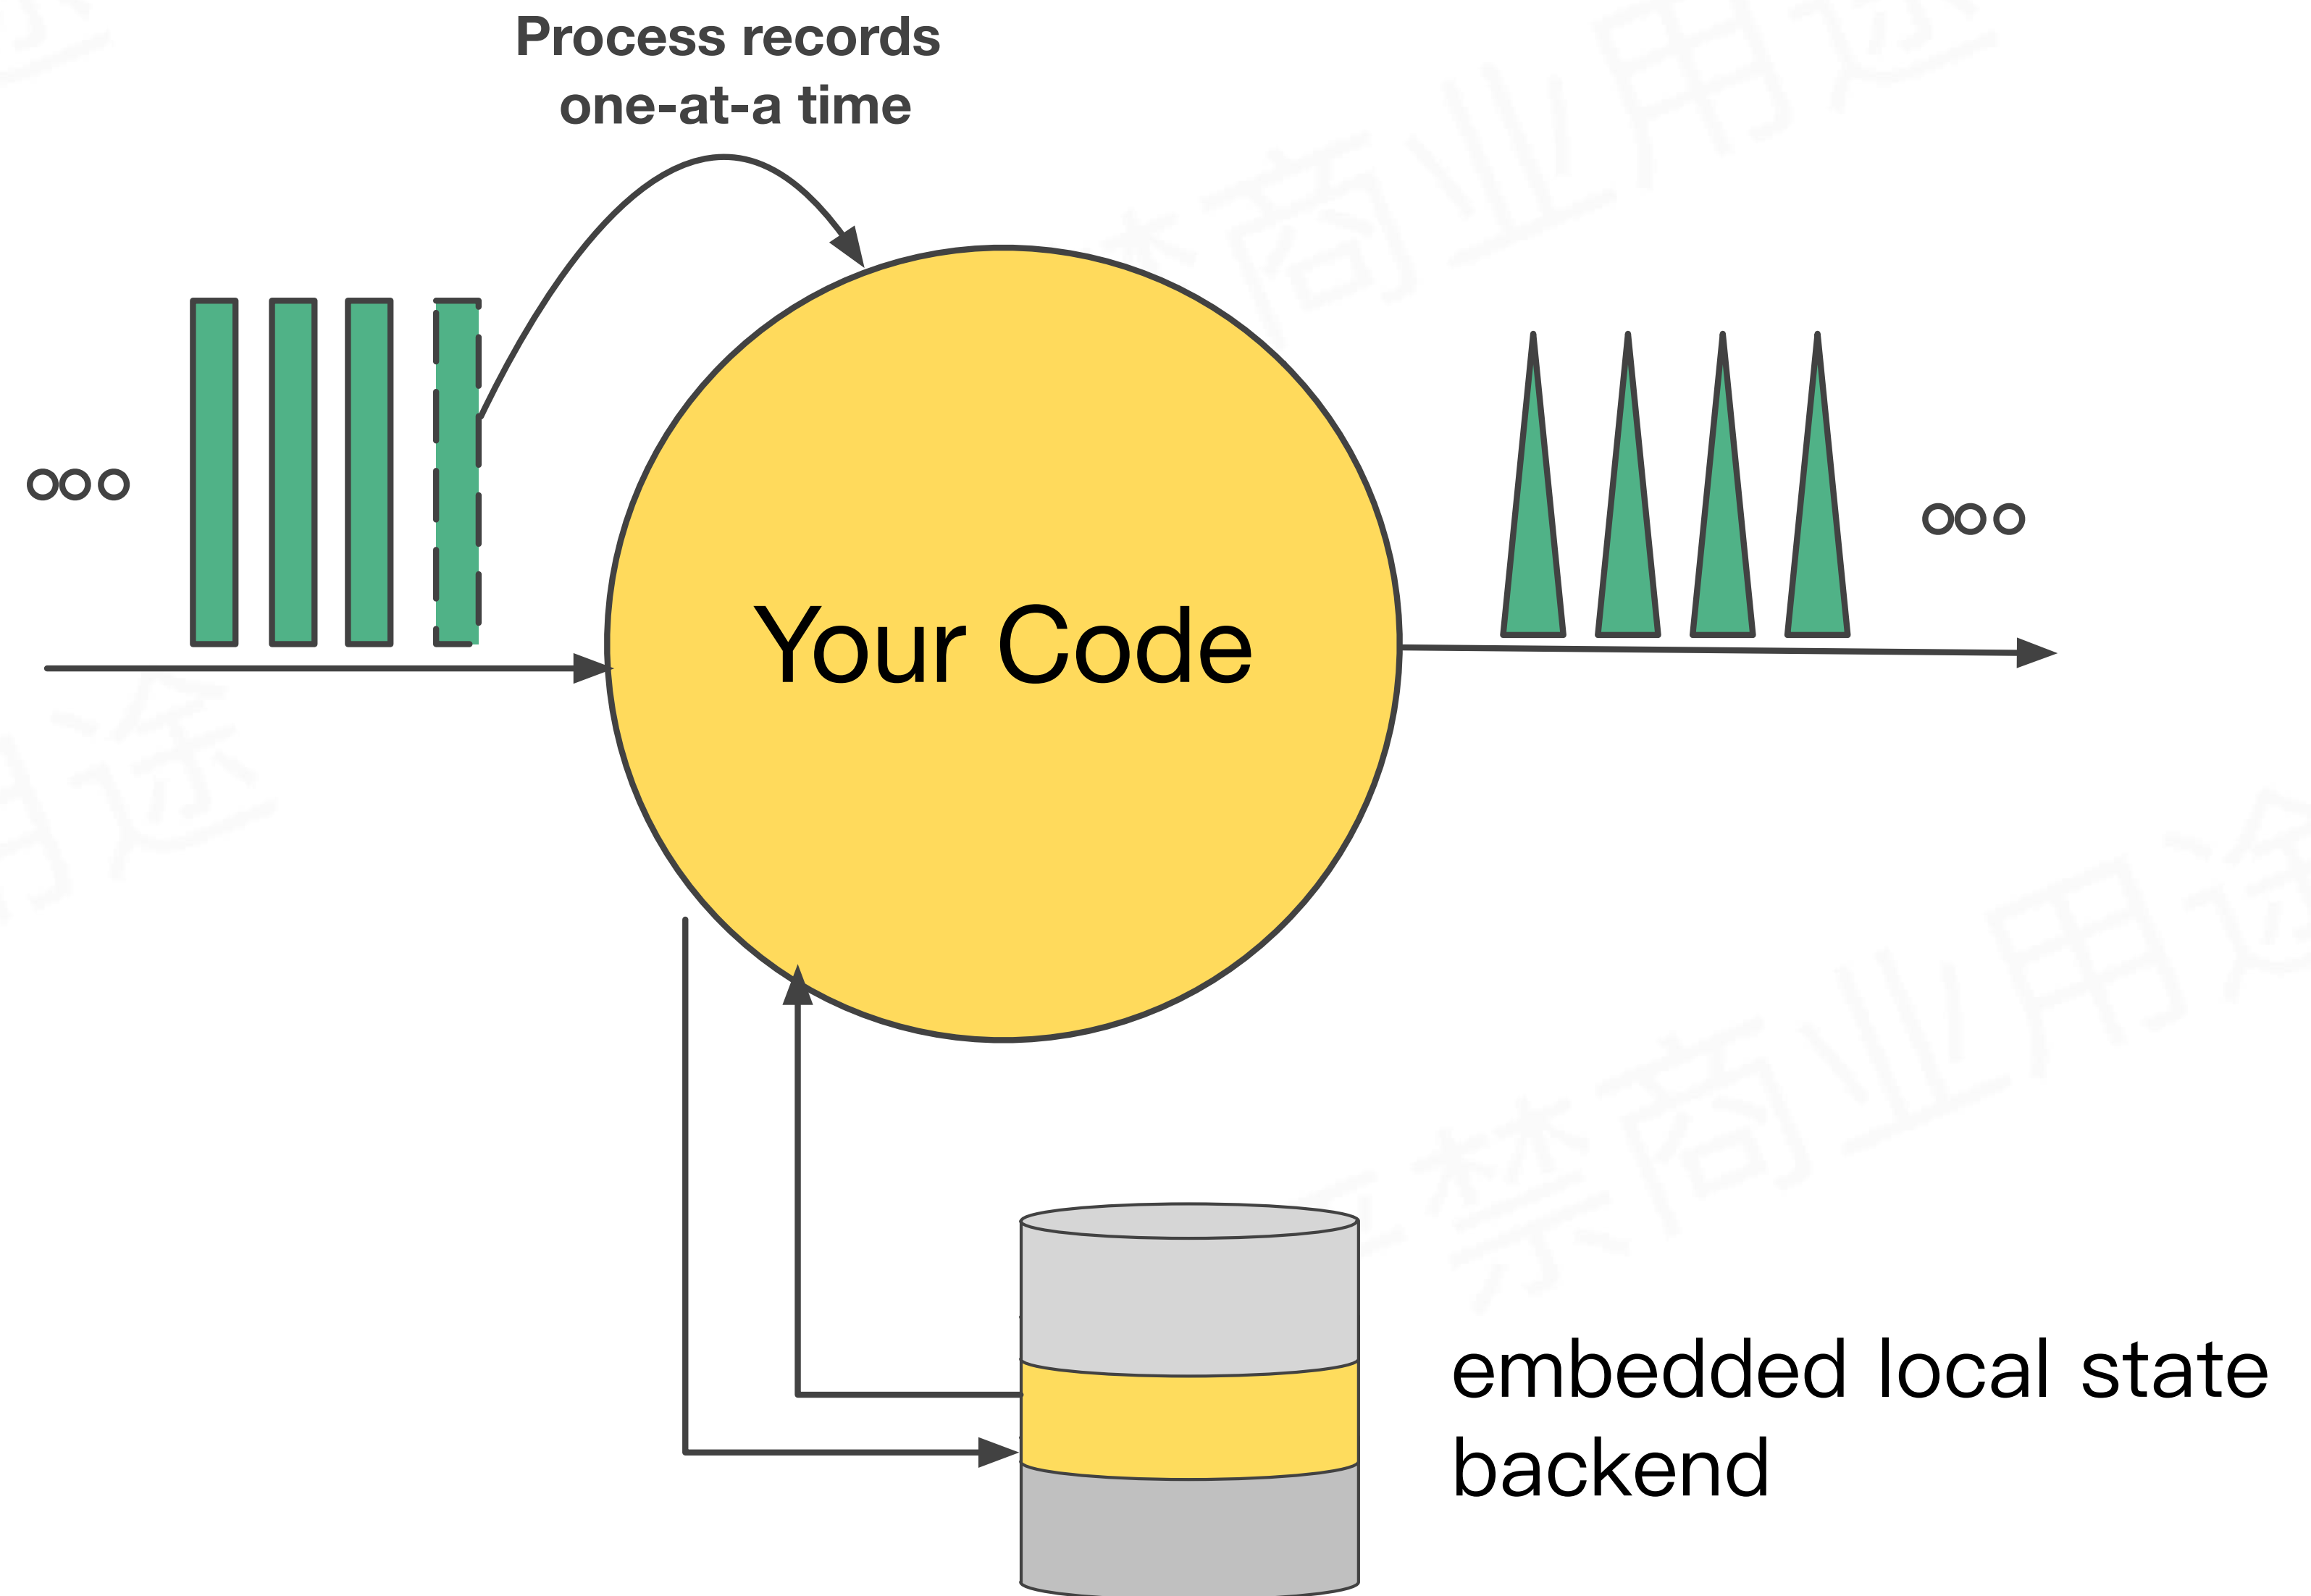
1. 执行上述代码
2. 本地启动netcat
3. 键入 `hello world` , 执行程序会输出什么?

```
$ nc -lk 9000
```

再次键入 `hello world` , 执行程序会输出什么?

# 什么是state

State:流式计算中持久化了的状态





## Keyed State

- 只能应用于 KeyedStream 的函数与操作中，例如Keyed UDF, window state
- Keyed State是已经分区/划分好的，每一个key只能属于某一个keyed state

# 什么是state **Keyed State**


再来看这  
段  
word  
count代码

```
env.socketTextStream("localhost",9000)
    // split up the lines in pairs (2-tuples) containing: (word,1)
    .flatMap(new Tokenizer())
    // group by the tuple field "0" and sum up tuple field "1"
    .keyBy(0).sum(1)
    .print();
```

创建  
KeyedStream  
(对key进行了划分,  
不同task上不会出  
现相同的key)

调用内置的  
StreamGroupedReduce UDF

```
env.socketTextStream("localhost",9000)
// split up the lines in pairs (2-tuples) containing: (word,1)
.flatMap(new Tokenizer())
// group by the tuple field "0" and sum up tuple field "1"
.keyBy(0).sum(1)
.print();
```



## Operator State

- 又称为non-keyed state，每一个operator state都仅与一个operator的实例绑定。
- 常见的operator state是source state，例如记录当前source的offset



# 什么是state

## Operator State

```
env.fromElements(WordCountData.WORDS)
    // split up the lines in pairs (2-tuples) containing: (word,1)
    .flatMap(new Tokenizer())
    // group by the tuple field "0" and sum up tuple field "1"
    .keyBy(0).sum(1)
    .print();
```

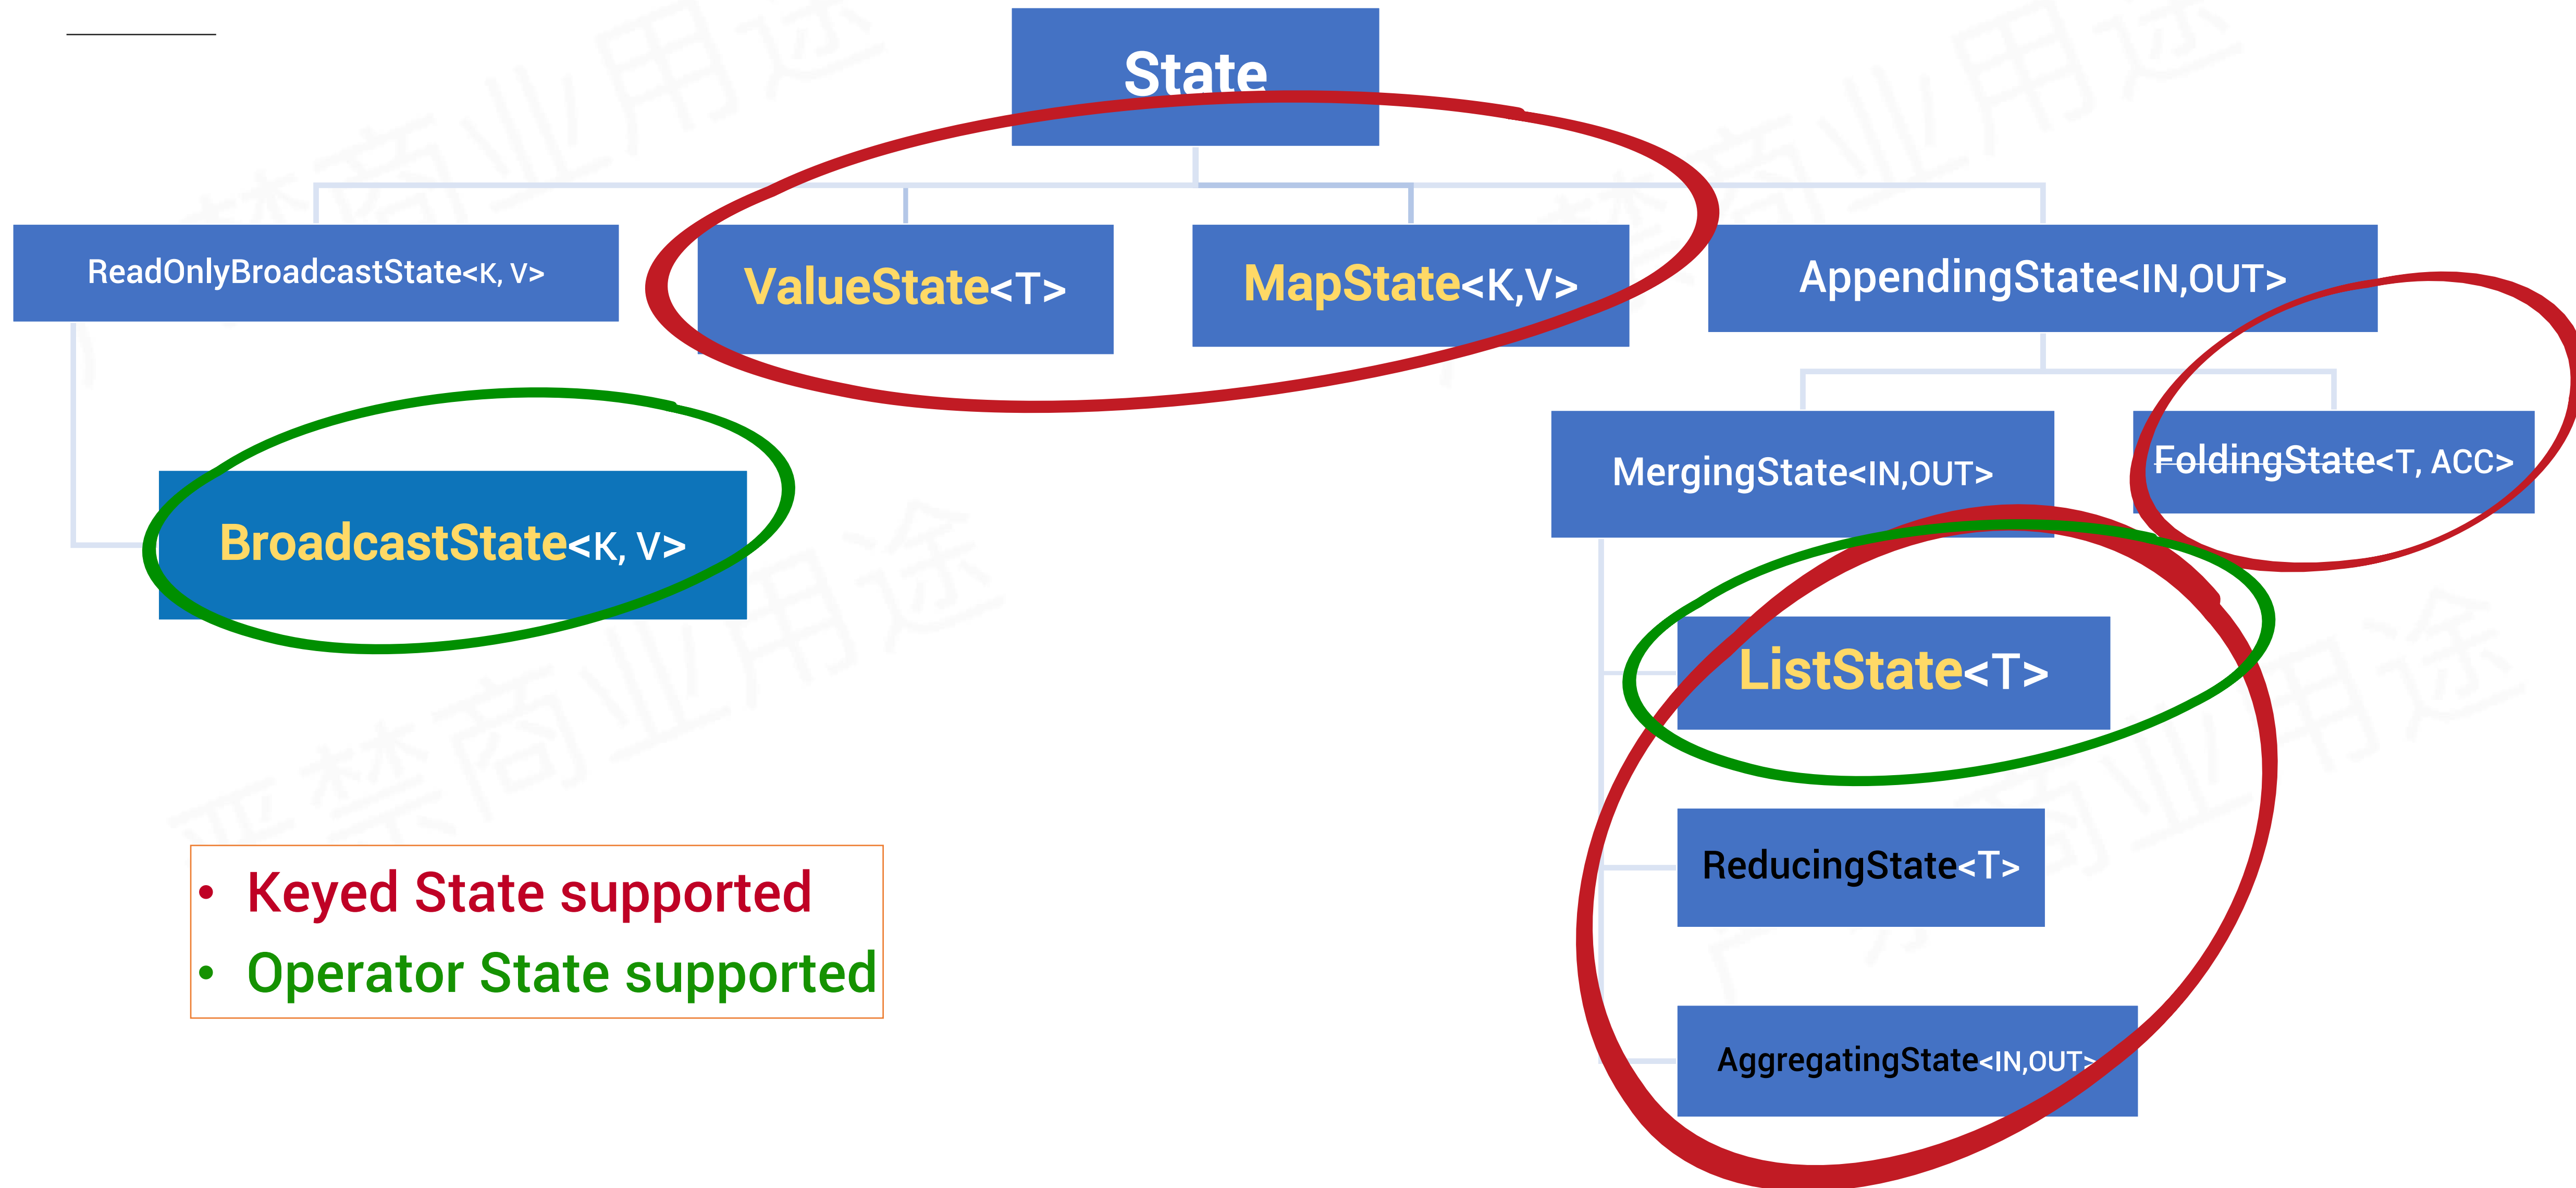
再来看一段  
word count代码

```
public class FromElementsFunction<T> implements SourceFunction<T>, CheckpointedFunction {
    private transient ListState<Integer> checkpointedState;

    public FromElementsFunction(TypeSerializer<T> serializer, T... elements) throws IOException {
        this(serializer, Arrays.asList(elements));
    }
}
```



# 什么是state



# 什么是state

- Managed State : 由Flink管理的state, 刚才举例的所有state均是managed
- Raw State : Flink仅提供stream可以进行存储数据, 对其而言只是一些bytes

```
public interface StateSnapshotContext extends FunctionSnapshotContext {  
    /**  
     * Returns an output stream for keyed state  
     */  
    KeyedStateCheckpointOutputStream getRawKeyedOperatorStateOutput() throws Exception;  
  
    /**  
     * Returns an output stream for operator state  
     */  
    OperatorStateCheckpointOutputStream getRawOperatorStateOutput() throws Exception;  
}
```

# 03

---

## 如何在Flink中使用state

使用指南

---

# 如何在Flink中使用state: **Keyed State**

再来看这段  
word count代码

```
env.fromElements(WordCountData.WORDS)
    // split up the lines in pairs (2-tuples) containing: (word,1)
    .flatMap(new Tokenizer())
    // group by the tuple field "0" and sum up tuple field "1"
    .keyBy(0).sum(1)
    .print();
```

创建KeyedStream  
(对key进行了划分,  
不同task上不会出现相  
同的key)

调用内置的StreamGroupedReduce





# 如何在Flink中使用state: Keyed State

```
public class StreamGroupedReduce<IN> extends AbstractUdfStreamOperator<IN, ReduceFunction<IN>>
    implements OneInputStreamOperator<IN, IN> {

    private transient ValueState<IN> values;

    @Override
    public void open() throws Exception {
        super.open();
        ValueStateDescriptor<IN> stateId = new ValueStateDescriptor<>(STATE_NAME, serializer);
        values = getRuntimeContext().getState(stateId);
    }

    @Override
    public void processElement(StreamRecord<IN> element) throws Exception {
        IN value = element.getValue();
        IN currentValue = values.value();

        if (currentValue != null) {
            IN reduced = userFunction.reduce(currentValue, value);
            values.update(reduced);
            output.collect(element.replace(reduced));
        } else {
            values.update(value);
            output.collect(element.replace(value));
        }
    }
}
```

通过 RuntimeContext 访问state

访问和修改当前key对应的state数值



# 如何在Flink中使用state: Operator State



Apache Flink

再来看这段  
word count代码

```
env.fromElements WordCountData.WORDS)
    // split up the lines in pairs (2-tuples) containing: (word,1)
    .flatMap(new Tokenizer())
    // group by the tuple field "0" and sum up tuple field "1"
    .keyBy(0).sum(1)
    .print();
```

调用内置的FromElementsFunction



# 如何在Flink中使用state: Operator State



Apache Flink

```
public class FromElementsFunction<T> implements SourceFunction<T>, CheckpointedFunction {

    private transient ListState<Integer> checkpointedState;

    @Override
    public void initializeState(FunctionInitializationContext context) throws Exception {
        Preconditions.checkState(this.checkpointedState == null,
            "The " + getClass().getSimpleName() + " has already been initialized.");

        this.checkpointedState = context.getOperatorStateStore().getListState(
            new ListStateDescriptor<>("from-elements-state", IntSerializer.INSTANCE)
        );

        if (context.isRestored()) {
            List<Integer> retrievedStates = new ArrayList<>();
            for (Integer entry : this.checkpointedState.get()) {
                retrievedStates.add(entry);
            }

            // given that the parallelism of the function is 1, we can only have 1 state
            Preconditions.checkArgument(retrievedStates.size() == 1, getClass().getSimpleName() + " retrieved invalid state.");

            this.numElementsToSkip = retrievedStates.get(0);
        }
    }

    @Override
    public void snapshotState(FunctionSnapshotContext context) throws Exception {
        Preconditions.checkState(this.checkpointedState != null,
            "The " + getClass().getSimpleName() + " has not been properly initialized.");

        this.checkpointedState.clear();
        this.checkpointedState.add(this.numElementsEmitted);
    }
}
```

通过 FunctionInitializationContext 访问state

在snapshotState时将状态数据存储在state中

# 04

---

## Checkpoint的执行机制

Checkpoint internal

---





*state-backend* 分类:

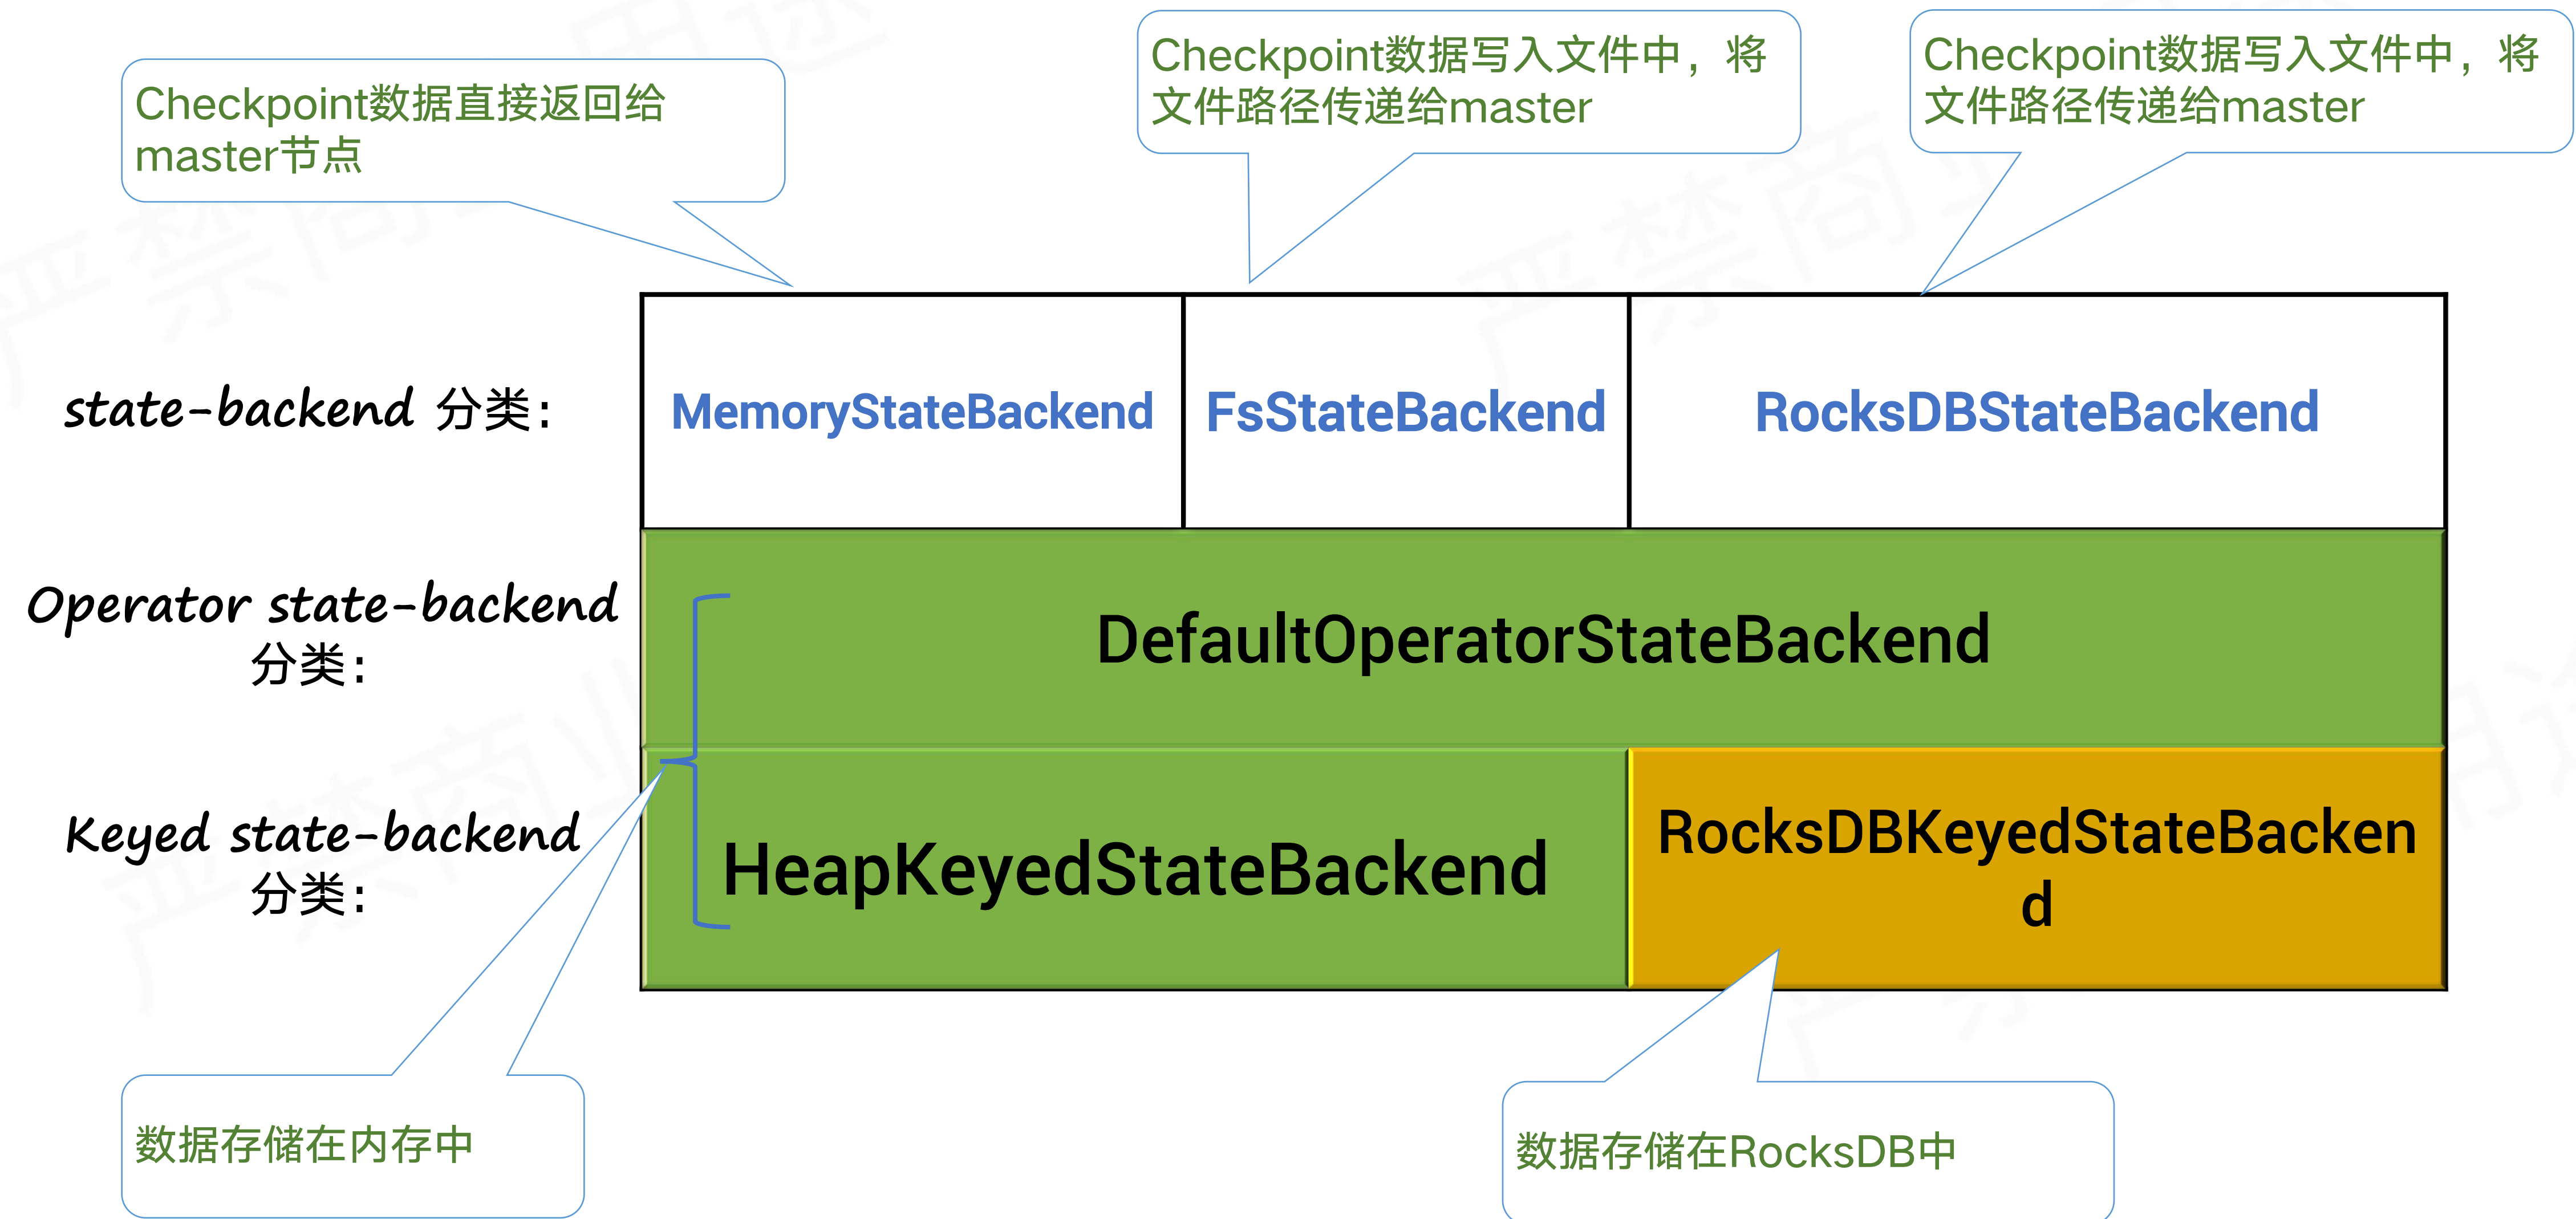
<b>MemoryStateBackend</b>	<b>FsStateBackend</b>	<b>RocksDBStateBackend</b>
<b>DefaultOperatorStateBackend</b>		
<b>HeapKeyedStateBackend</b>		<b>RocksDBKeyedStateBackend</b>

*Operator state-backend*  
分类:

*Keyed state-backend*  
分类:

```
StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();  
env.setStateBackend(new FsStateBackend("hdfs://namenode:40010/flink/checkpoints"));
```

# State的存储





# State的存储

## HeapKeyedStateBackend 存储格式


- 支持异步checkpoint（默认）：`CopyOnWriteStateTable<K, N, S>[]`，整体相当于一个map
- 仅支持同步checkpoint：`Map<N, Map<K, S>>[]`，由嵌套map的数组构成
- 在MemoryStateBackend内使用时，checkpoint序列化数据阶段默认有最大5MB数据的限制

## RocksDBKeyedStateBackend 存储格式

每个state都存储在一个单独的column family内  
keyGroup, Key和Namespace进行序列化存储在DB作为key

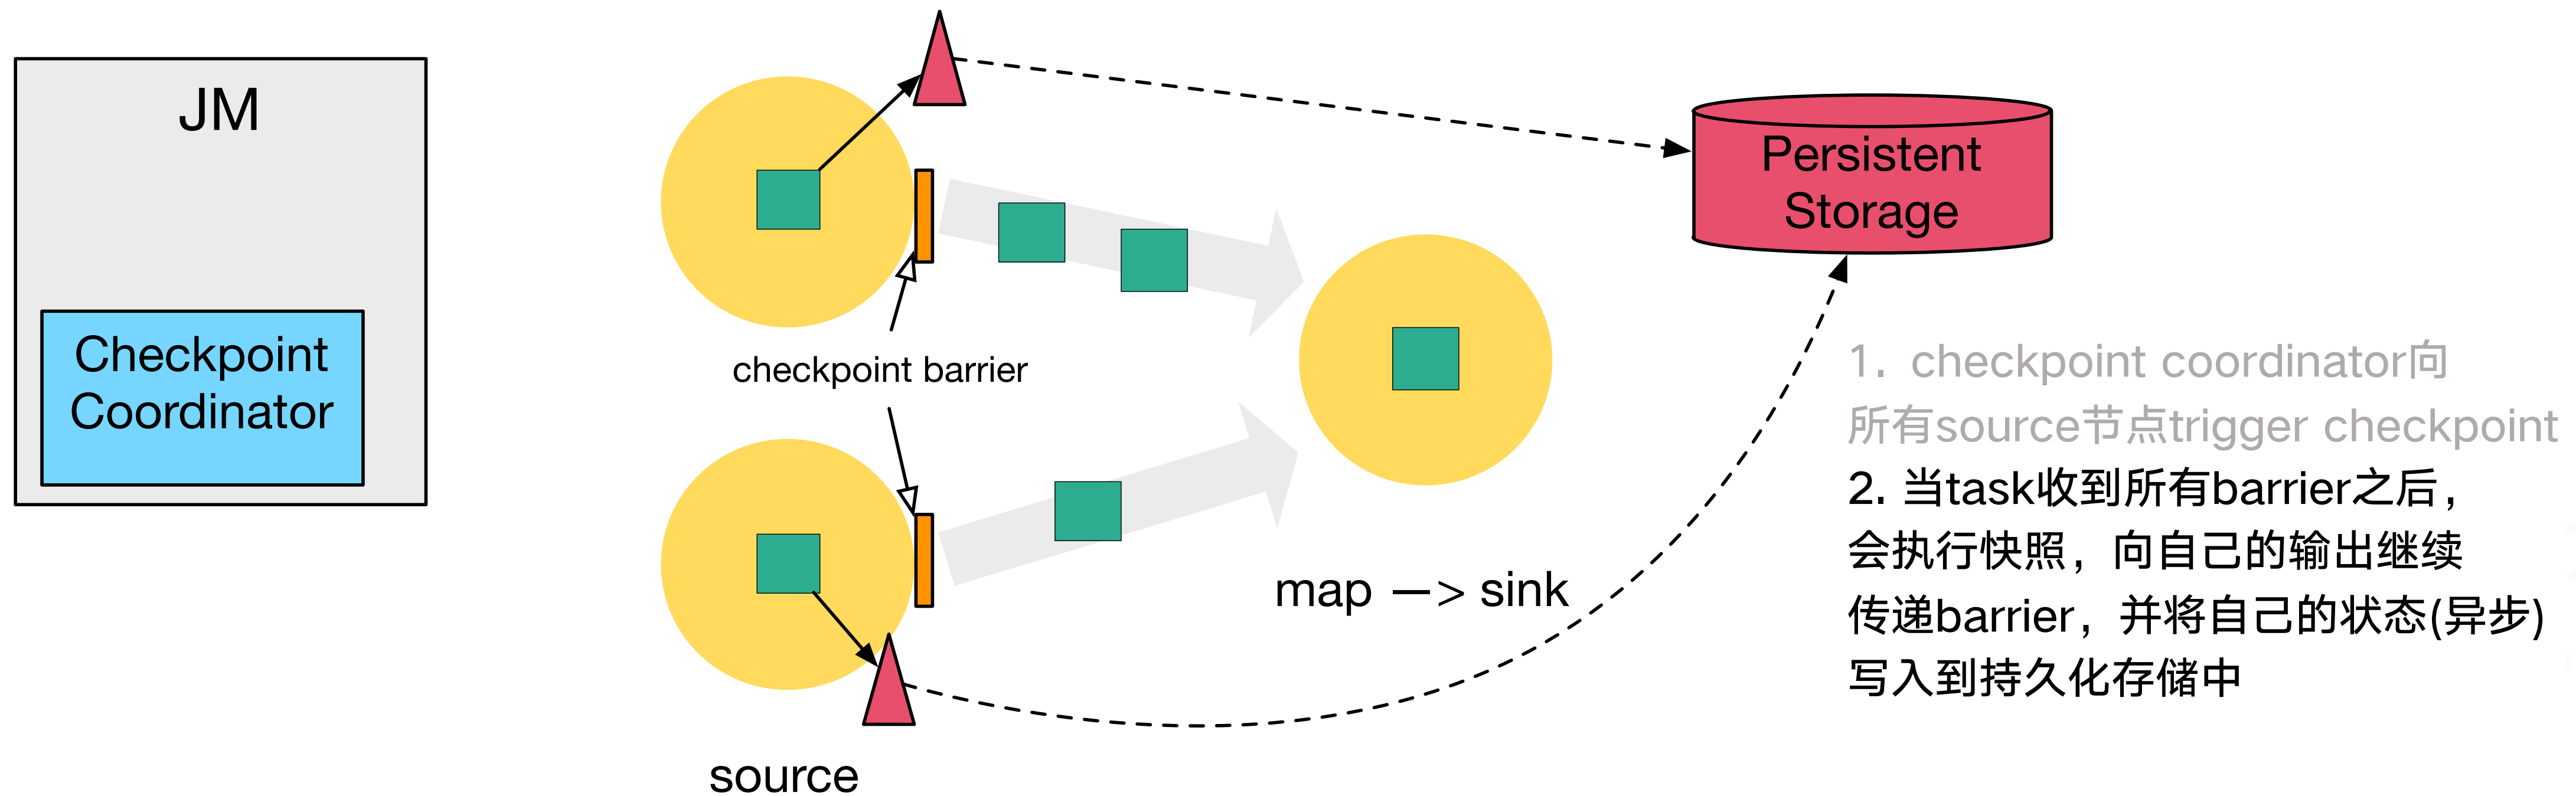
State1		State2	
KeyGroup + Key + Namespace	value	KeyGroup + Key + Namespace	value
(1, K1, Window(10, 20))	v1	(2, K2, Window(10, 20))	v2
(1, K3, Window(10, 20))	v3	(2, K4, Window(10, 25))	v4
...	...	...	...

# Checkpoint执行流程




1. checkpoint coordinator  
向所有source节点trigger checkpoint

# Checkpoint执行流程




# Checkpoint执行流程




1. checkpoint coordinator向所有source节点trigger checkpoint
2. 当task收到所有barrier之后，会执行快照，向自己的输出继续传递barrier，并将自己的状态(异步)写入到持久化存储中
3. 当task完成备份后，会将备份数据的地址（state handle）通知给checkpoint coordinator



# Checkpoint执行流程




# Checkpoint执行流程




1. checkpoint coordinator向所有source节点trigger checkpoint
2. 当task收到所有barrier之后，会执行快照，向自己的输出继续传递barrier，并将自己的状态(异步)写入到持久化存储中
3. 当task完成备份后，会将备份数据的地址（state handle）通知给checkpoint coordinator

# Checkpoint执行流程



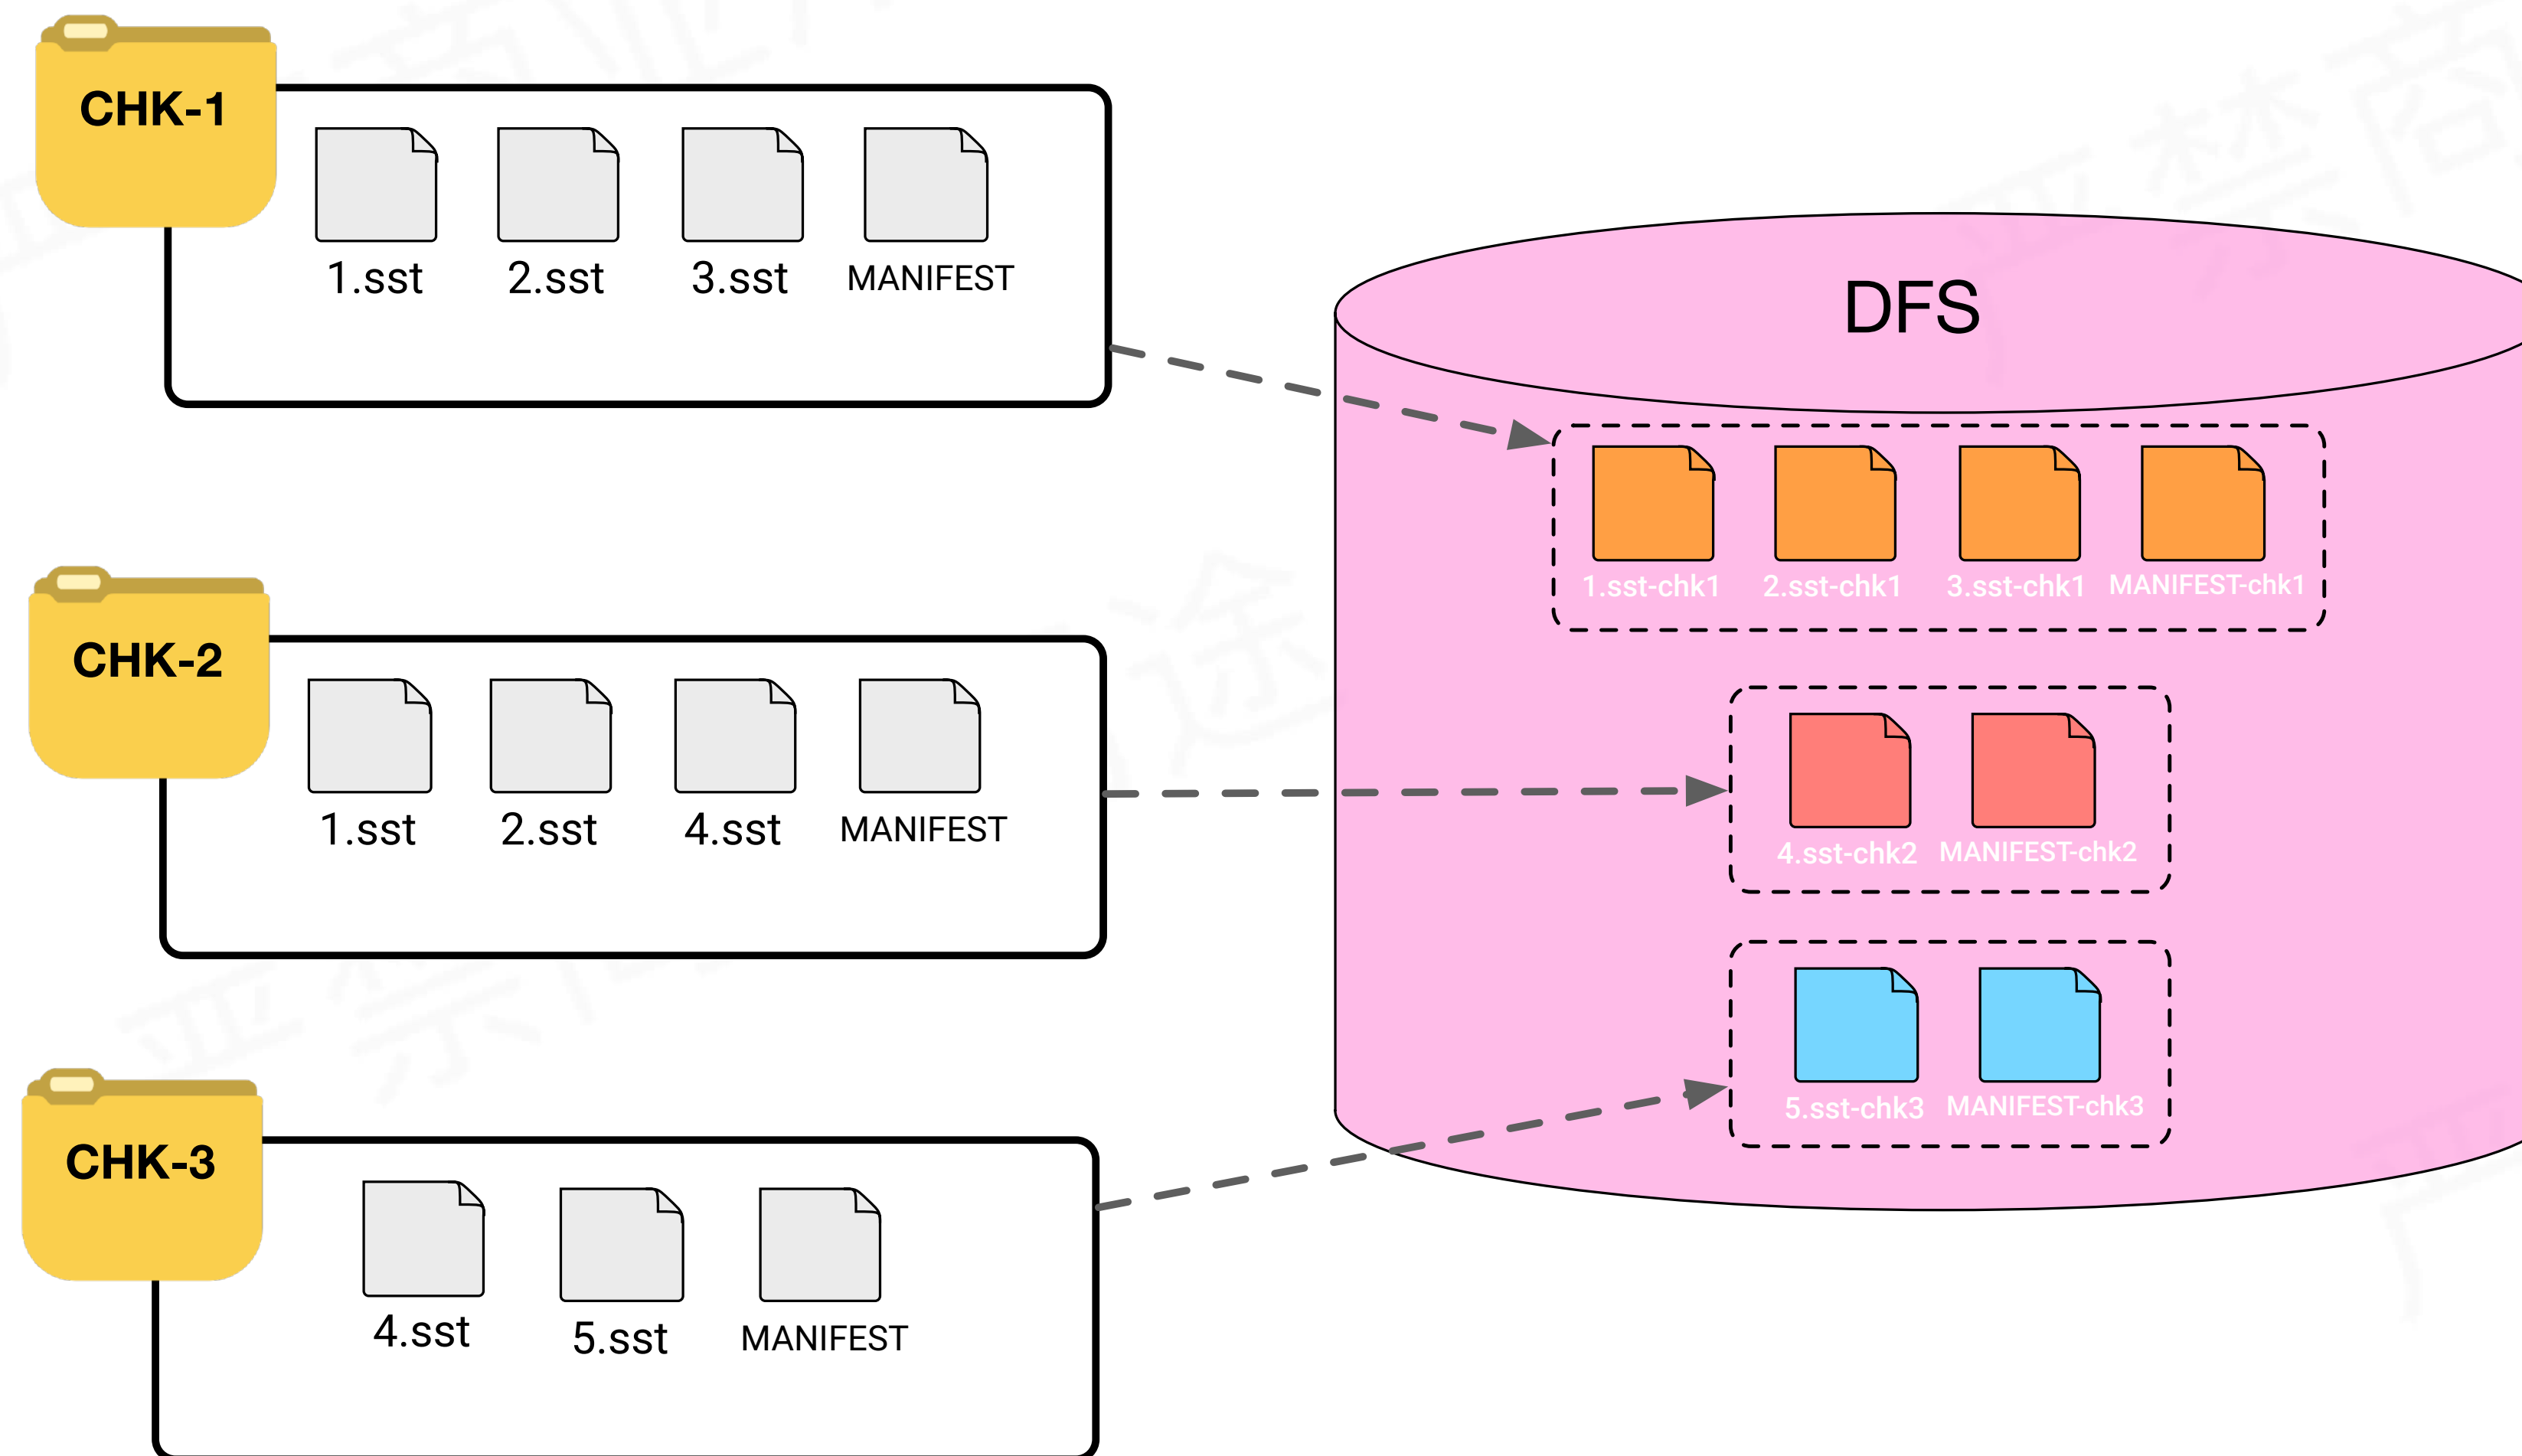
# Checkpoint执行流程



- 为了实现EXACTLY ONCE语义，Flink通过一个input buffer将在对齐阶段收到的数据缓存起来，等对齐完成之后再进行处理。
- 对于AT LEAST ONCE语义，无需缓存收集到的数据，会对后续直接处理，所以导致restore时，数据可能会被多次处理。
- Flink的checkpoint机制只能保证Flink的计算过程可以做到EXACTLY ONCE，end-to-end的EXACTLY ONCE需要source和sink支持



# 基于RocksDB的增量checkpoint



1. 本地snapshot目录创建当前DB内容的备份
2. 与上一次成功的checkpoint本地sst文件列表比对，将不在其中的文件上传到DFS中
3. 所有文件都会重命名防止冲突
4. 包含了所有新旧文件的handle返回给checkpoint coordinator



# 从已停止的作业进行状态恢复

---

## Savepoint

用户通过命令触发，由用户管理其创建与删除

标准化格式存储，允许作业升级或者配置变更


用户在恢复时需要提供用于恢复作业状态的savepoint路径

## Externalized Checkpoint

Checkpoint完成时，在用户给定的外部持久化存储保存

当作业FAILED（或者CANCELED）时，外部存储的checkpoint会保留下来


用户在恢复时需要提供用于恢复的作业状态的checkpoint路径



© Apache Flink Community China 严禁商业用途


# 从已停止的作业进行状态恢复 Operator State的改并发

ListState




所有task的元素均匀划分  
给新的task  
(Even-split redistribution)

UnionListState



所有task的元素全部划分给新  
的task  
(Union redistribution)

BroadcastState <K, V>



所有task的state相同，  
改并发时，新的task获  
得state的一个备份





Apache Flink

THANKS

Flink China社区大群



扫一扫群二维码，立刻加入该群。